



Cloud Computing et informatique distribuée

Cours : Hadoop

Francesca Bugiotti

CentraleSupélec

May 14, 2020



Objectives

- List the main components of Hadoop ecosystem and their function
- Understand the usage of HDFS and MapReduce in Hadoop
- Explain Spark architecture
- Summarize how Spark manages and executes code on Clusters



Plan

1 Hadoop



Hadoop [2]

Apache Hadoop is an open-source software framework for distributed storage and distributed processing of very large data sets on computer clusters built from commodity hardware

Characteristics

Hadoop components provide **scalability** to store and process large volumes of data on **commodity hardware**

Most of Hadoop components provide the possibility to **recover** from failures

It handles **different data types**



Hadoop

History

- 2004 - Google published a paper about their in-house processing framework called MapReduce [1]
- 2005 - Yahoo released an open-source implementation based on this framework called Hadoop

Now: more than 100 open-source projects on Hadoop



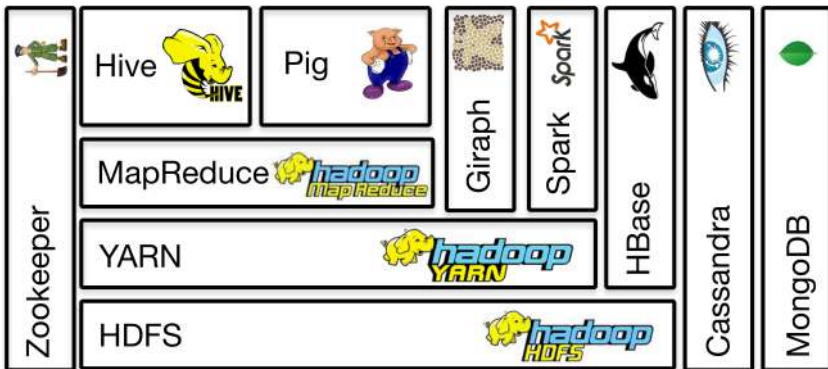
Hadoop

Many projects and components but on the **core part** we can identify:

- Hadoop Distributed File System (HDFS)
- Hadoop YARN
- Hadoop MapReduce



Hadoop





Hadoop

- **HDFS**: distributed file system on commodity hardware, scalable and reliable storage
- **YARN**: flexible scheduling and resource management over HDFS storage
- **MapReduce**: programming model for parallel computing
- **Pig**: augments data modeling of MapReduce with flow modeling
- **Hive**: augments data modeling of MapReduce with relational Algebra
- **Giraph**: processes of large scale graphs
- **Storm**, **Spark**, and **Flink** real time and in memory processing of Big Data
- **NoSQL projects** like Cassandra (created by Facebook) that also uses HBase
- **Zookeeper**: centralized management system for synchronization, configuration, and high availability



References I



Jeffrey Dean and Sanjay Ghemawat.

Mapreduce: Simplified data processing on large clusters.

In *6th Symposium on Operating System Design and Implementation (OSDI 2004)*, San Francisco, California, USA, December 6-8, 2004, pages 137–150, 2004.



Tom White.

Hadoop: The Definitive Guide.

O'Reilly Media, Inc., 1st edition, 2009.



Cloud Computing et informatique distribuée

Cours : HDFS

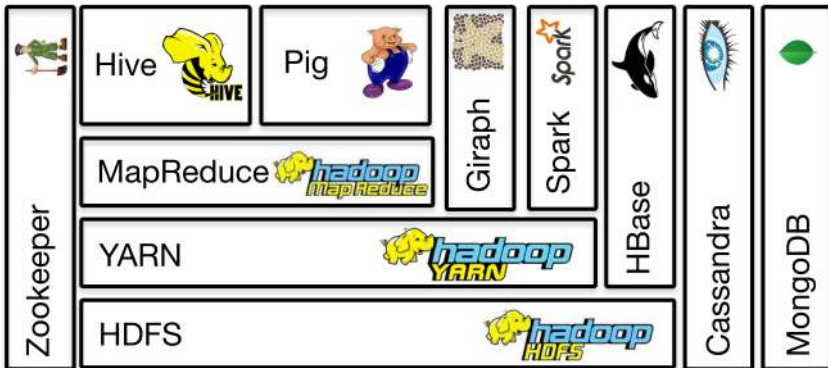
Francesca Bugiotti

CentraleSupélec

May 14, 2020



Hadoop






Plan

1 Hadoop HDFS



HDFS



Hadoop Distributed File System

A storage layer for Big Data

- Scalability
- Reliability



HDFS

Some numbers

- Hortonworks: 200 petabytes
- Single cluster of 4.500 servers
- A billion of files and blocks

You do not run out of space:

- You add more nodes to increase the space

Scalability:

- Partition and splitting
 - Parallel access



HDFS

More numbers

- Gigabytes to terabyte
- Chunk size (the size of each piece) 64 megabytes

Fault tolerance

- Replica of file blocks
- By default 3 copies (configurable) -> replication factor

Many data types:

- Input file/output file format
- Flexibility:
 - Geospatial data can be read as vectors or rasters



HDFS

Two components

Name Node:

- Responsible for metadata
 - Name, location in the directory hierarchy, etc.
- Decides where to store data
- There is one in each cluster

Data Nodes:

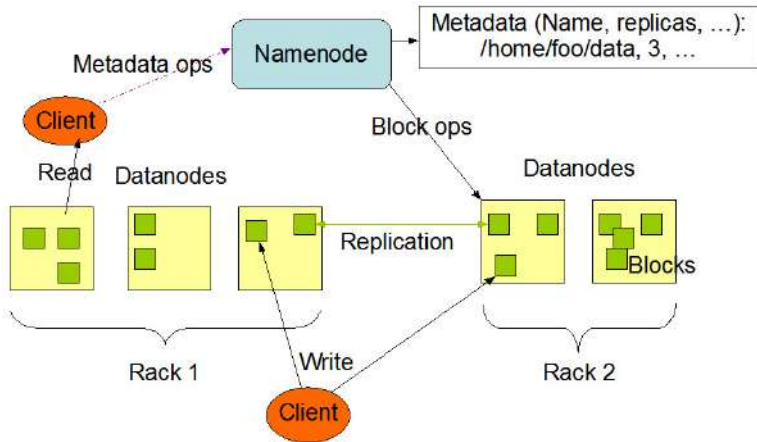
- Block storage
- Does the block creation, deletion, and replication
- There is one in each node

Master/slave relationship



HDFS [?]

HDFS Architecture





HDFS in practice

- Copy myfile.txt into HDFS:

```
hadoop fs -copyFromLocal myfile.txt
```

- Verify the existing files in HDFS:

```
hadoop fs -ls
```

- Delete a file:

```
hadoop fs -rm myfile.txt
```

Cloud Computing et informatique distribuée

Cours : Yarn

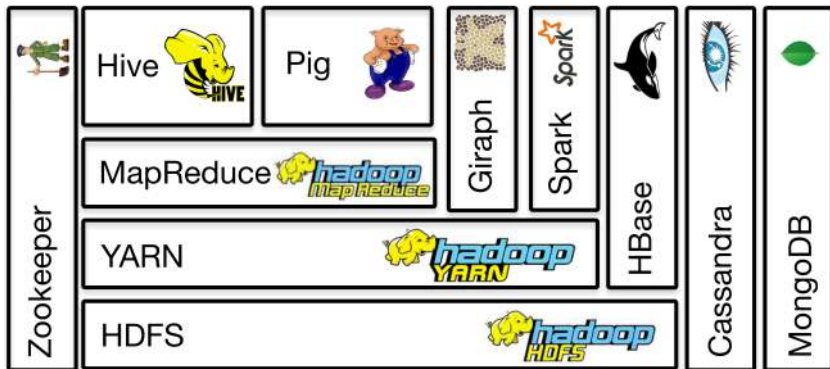
Francesca Bugiotti

CentraleSupélec

May 14, 2020



Hadoop






Plan

1 YARN

YARN [1]



Hadoop Resource Manager

Interacts with applications and schedules resources for their usage

- Multiple applications
- Beyond MapReduce
- Beyond the data parallel programming model



YARN

Added in Hadoop Ecosystem in a second moment

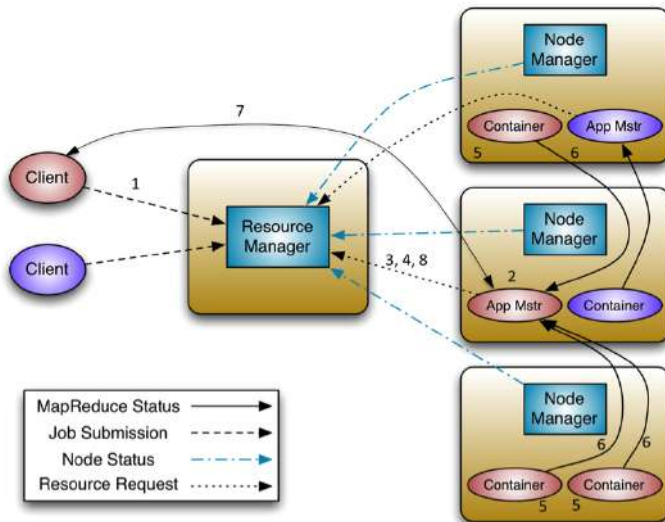
- Allows the support for not MapReduce applications
- Enables graph analytics
- Streaming data

Main components:

- Resource Manager
- Node Manager
- Container



YARN





YARN

- The **resource manager** controls all the resources, and decides which one gets what
- **Node manager** operates at machine level and is in charge of a single machine

Application level:

- Each application gets an **application master**
 - Resources negotiation with the Resource Manager
 - Interaction with Node Manager to get the tasks completed
- **Container**
 - A collection of CPU memory disk network and other resources within the compute node



YARN

- Increases CPU utilization
- 1000 machines to their 2500 machines cluster
- Twice the number of jobs run before

Many distributed applications over the same Hadoop cluster



References I



Hadoop HDFS.

`http:`

`//hadoop.apache.org/docs/r1.2.1/hdfs_design.html.`

Accessed 2016.



Cloud Computing et informatique distribuée

Cours : Hadoop-MapReduce

Francesca Bugiotti

CentraleSupélec

May 14, 2020



Plan

1 Hadoop MapReduce



Hadoop MapReduce


With MapReduce you must create Map and Reduce tasks and run them

- MapReduce job usually splits the input data-set into independent chunks
 - the process works in parallel manner
- The framework sorts and shuffles the outputs of the maps
- The output of the map is given in input to the reduce tasks
- Both the input and the output of the job are stored in a file-system

The framework takes care of **scheduling** tasks, **monitoring** them and **re-executes** the failed tasks.



MapReduce





MapReduce problems

- **Programming model**
 - Hard to implement everything as a MapReduce program
 - Iterative algorithms
 - Machine Learning, Graphs & Network Analysis
 - Interactive Data Mining
 - R, Excel-like computations, ad hoc reporting, etc.
 - Multiple MapReduce steps for simple operations
 - Lack of control structures and data types



MapReduce problems

- **Efficiency**
 - High communication cost
 - Frequent writing of output to disk
 - Limited exploitation of main memory
- **Real-time processing**
 - A MapReduce job requires to scan the entire input
 - Stream processing and random access impossible
- **Programming language**
 - Native support for Java only



Solutions?

Leverage to **memory**:

- Replace disks with **SSD** (Solid State Drive)
- Load data into memory



Cloud Computing et informatique distribuée

Cours : Hadoop-Spark

Francesca Bugiotti

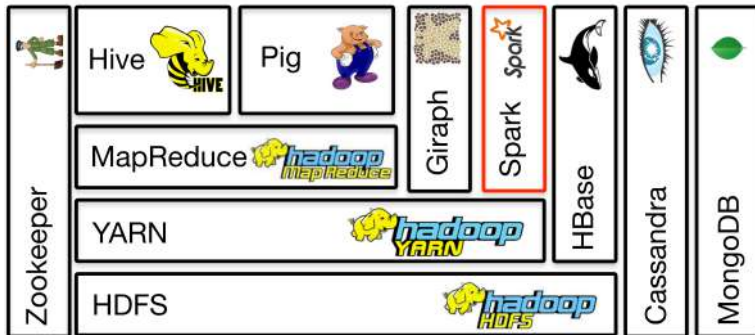
CentraleSupélec

May 14, 2020




Plan

1 Spark



History

- Spark [1, 2] project started in 2009
- Developed originally at UC Berkeley's AMPLab
- Open sourced in 2010
- Shark started summer 2011, alpha April 2012
- It is now the most popular project for big data analysis [3]





Spark

- Separate, fast, MapReduce-like engine
 - In-memory data storage for very fast iterative queries
 - General execution graphs and powerful optimizations
 - Up to 40x faster than Hadoop
- Compatible with Hadoop's storage APIs
 - Can run on top of an Hadoop cluster
 - Can read/write to any Hadoop-supported system, including HDFS, HBase, SequenceFiles, etc.
- Not a modified version of Hadoop



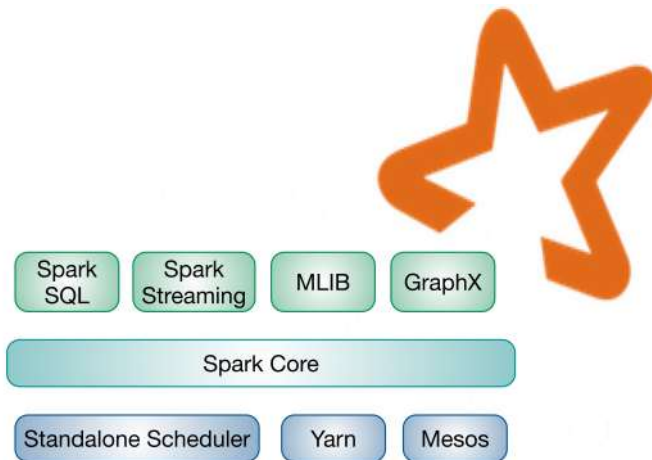
Spark

MapReduce greatly simplified big data analysis, but as soon as it got popular, users wanted more:

- Complex, multi-stage applications
 - Iterative graph algorithms and machine learning
 - RDD (Resilient Distributed Datasets) as unit of manipulation
- Efficiency
- Interactive ad-hoc queries
- Both multi-stage and interactive apps require faster data sharing across parallel jobs

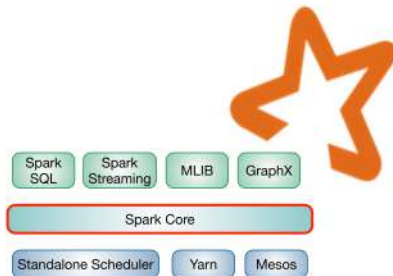


Spark Stack





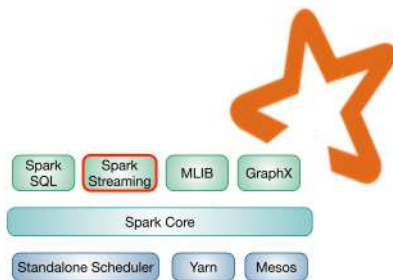
Spark Core



- Memory management
- Distributed scheduling
- Fault tolerance
- API for distributed datasets



Spark Streaming

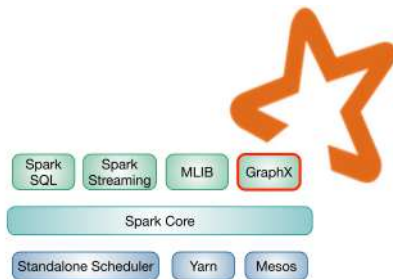


It implements an incremental stream processing using a model called “discretized streams”

- Input is split in small batches
- Regular combinations with states stored into RDD (Resilient Distributed Datasets)



Spark GraphX

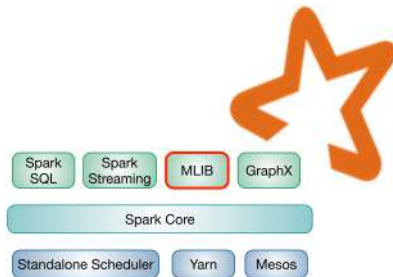


A graph computation interface that treats each graph as a directed multigraph with properties attached to each vertex and edge

- Variety of graph algorithms
- Flexible, fault tolerant, intuitive



MLib

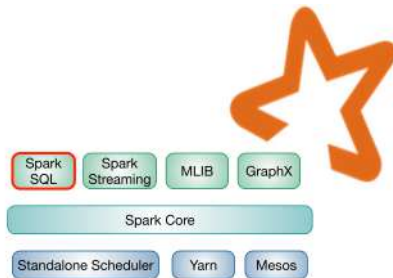


Machine learning library

- More than 50 algorithms implemented

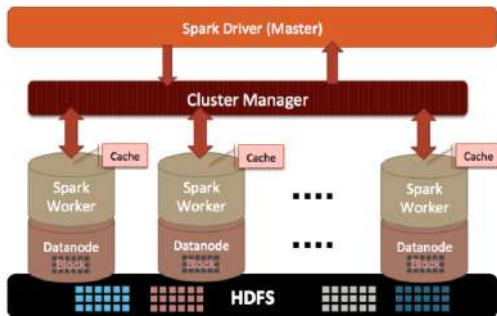


Spark SQL




SQL queries on Spark with techniques similar to analytical databases

- Column storage
- Cost-based optimization
- Code generation for query execution






Cluster Architecture [4]




Cluster Architecture [4]




- Applications run as **independent sets of processes** on a cluster
- Applications are coordinated by the **SparkContext** object in the main program (the driver program)
- Each application gets its own **executor processes**
- The execution process stays up for the duration of the whole application and runs tasks in **multiple threads**

Cluster Architecture [4]



- To run on a cluster, the SparkContext can connect to **several types** of cluster managers, which allocate resources across applications
- Once connected, Spark acquires executors on nodes in the cluster, which run computations and store data
- Next, it sends the application code to the executors
- Finally, SparkContext sends tasks to the executors to run

Cluster Architecture [4]



Two main components:

- A driver program
- A set of worker nodes



Driver Program

Where the application starts:

- Distributes RDDs on a computational cluster
- Creates a connection to a Spark cluster through a `SparkContext` object
- The default `SparkContext` in the Spark shell is an object (called `SC` for `SparkContext`)
- Manages a potentially large number of nodes called `worker nodes`

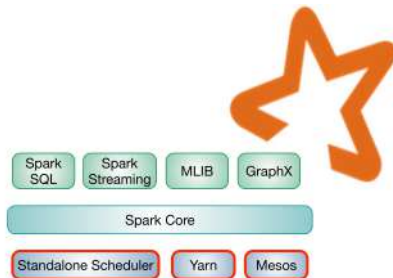


Worker Node

- Keeps a running Java virtual machine, called the **executor**
 - The core for application execution
 - Executors can execute tasks related to mapping stages or reducing stages or other Spark specific pipelines
- In a real Big Data scenario, we have many worker nodes running tasks internally
- Automatic provisioning and restarting of these nodes handled by the cluster manager




Worker Node



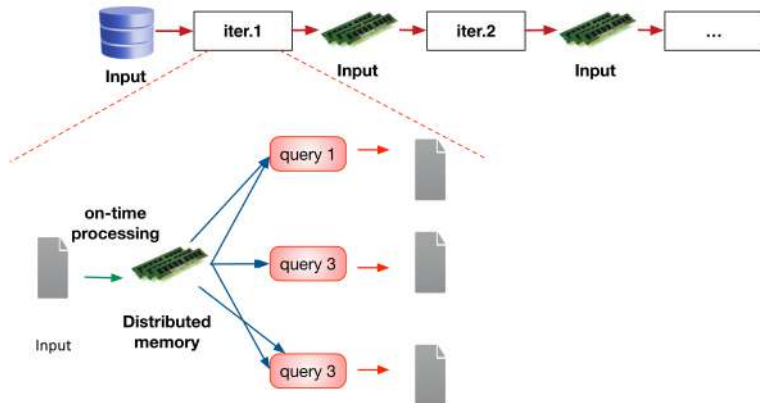
Three interfaces for cluster management:

- Spark's standalone cluster manager
- The Apache Mesos
- Hadoop YARN


Data in MapReduce



Data in Spark



Data in Spark



- Extract a working set
- Cache the working set
- Query the working set repeatedly



RDD: Resilient Distributed Dataset

- Immutable Data Structure
- In-memory
- Fault tolerant
- Parallel data structure
- Controlled partitioning
- Manipulated using a rich set of operators



RDD: Resilient Distributed Dataset

- Distributed collection of objects that can be cached in memory across cluster nodes
- Manipulated through various parallel operators
- Automatically rebuilt on failure
- Can persist in memory, on disk, or both
- Can be partitioned to control parallel processing
- Interface:
 - Clean **language-integrated** API for Scala, Python, and Java
 - Can be used interactively from **Scala console**



RDD: Resilient Distributed Dataset

Currently two types of RDDs

- **Parallelized collections:** created by executing operators on an existing programming collection
 - Developer can specify the number of slices to cut the dataset into
 - Ideally 2-3 slices per CPU
- **Hadoop Datasets:** created from any file stored on HDFS or other storage systems supported by Hadoop (S3, Hbase, etc.)
 - These are created using SparkContext's `textFile` operator
 - Default number of slices in this case is 1 slice per file block



Operators over RDD

Programmer can perform three types of operations:

- 1 Transformations
- 2 Actions
- 3 Persistence



Operators over RDD

1 Transformations:

- Create a new dataset from an existing one
- Lazy in nature
- They are executed only when some **action** is performed

Example

- `Map(func)`
- `Filter(func)`
- `Distinct()`



Transformations

- Transformations lazy operations on a RDD
- They return RDD objects or collections of RDD
- Are not executed immediately, but only after an action has been executed

Two kind of transformations:

- Wide transformations
- Narrow transformations



Narrow Transformations

- They are the result of `map`, `filter` and such that is from the data from a **single partition only**
 - i.e. it is self-sustained
- An output RDD has partitions with records that originate from a single partition in the parent RDD
- Spark groups narrow transformations as a **stage**



Wide Transformations

- They are the result of `groupByKey` and `reduceByKey`
- The data required to compute the records in a single partition may reside in many partitions of the parent RDD
- All of the tuples with the same key must end up in the same partition, processed by the same task
- Spark must execute **RDD shuffle**, which transfers data across cluster and results in a **new stage** with a **new set of partitions**



Transformations

- `filter`
- `map`
- `flatMap`
- `sample`
- `union`
- `intersection`
- `distinct`
- `groupByKey`
- `reduceByKey`
- `sortByKey`
- `join`
- `cogroup`
- `cartesian`



Operators over RDD

2 Actions:

- Returns a **value** to the driver program
- Exports data to a storage system after performing a computation

Example

- `Count()`
- `Reduce(func)`
- `Collect`
- `Take()`



Actions

- Actions are operations that return values
 - i.e. any RDD operation that returns a value of any type but an RDD is an action
 - e.g., `saveAs`, `collect`, `take`, `reduce`, etc.
- Actions are **synchronous**: they trigger execution of RDD transformations to return values
- Until no action is fired, the data to be processed is not even accessed
- Only actions can materialize the entire process with real data
- Cause data to be returned to driver or saved to output
- Cause data retrieval and execution of all transformations on RDDs



Actions

- `reduce`
- `collect`
- `count`
- `first`
- `take`
- `takeSample`
- `takeOrdered`
- `saveAsTextFile`
- `saveAsSequenceFile`
- `foreach`



Operators over RDD

3 Persistence:

- For caching datasets in-memory for future operations
- Option to store on disk or RAM or mixed (Storage Level)

Example

- `Persist()`
- `Cache()`



Persistence

By default, each transformed RDD is recomputed each time you run an action on it, unless you specify the RDD to be cached in memory

- RDD can be persisted on disks as well
- Caching is the key tool for iterative algorithms
- Using `persist()`, one can specify the **Storage Level** for persisting an RDD
- `cache()` is just a short hand for default storage level, which is **MEMORY_ONLY**



Persistence

Storage Level for `persist()`:

- `MEMORY_ONLY`
- `MEMORY_AND_DISK`
- `DISK_ONLY`
- `MEMORY_ONLY_2`, `MEMORY_AND_DISK_2`, etc.

Which Storage level is best?

- Try to keep in-memory as much as possible
- Try not to spill to disc unless your computed datasets are memory expensive
- Use replication only if you want fault tolerance



How does Spark work?

- User applications **create** RDDs, **transform** them, and **run** actions
- This results in a DAG (Directed Acyclic Graph) of **operators**
- DAG is compiled into **stages**
- Each stage is executed as a series of **Tasks** (one Task for each Partition)
- **Actions** drive the execution




Example in Python

```
text_file = sc.textFile(''hdfs://... '')
counts = text_file.flatMap(lambda line: line.split(" "))
                    .map(lambda word: (word, 1))
                    .reduceByKey(lambda a, b: a + b)
output = counts.collect()
output.saveAsTextFile(''hdfs://... '')
```


Example in Python

```
text_file = sc.textFile(''hdfs://... '')
```




Example in Python

```
text_file = sc.textFile('hdfs://...')  
counts = text_file.flatMap(lambda line: line.split(" "))
```




Example in Python

```
text_file = sc.textFile('hdfs://...')  
counts = text_file.flatMap(lambda line: line.split(" "))  
                  .map(lambda word: (word, 1))
```




Example in Python

```
text_file = sc.textFile('hdfs://...')  
counts = text_file.flatMap(lambda line: line.split(" "))  
                    .map(lambda word: (word, 1))  
                    .reduceByKey(lambda a, b: a + b)
```




Example in Python

```
text_file = sc.textFile('hdfs://...')  
counts = text_file.flatMap(lambda line: line.split(" "))  
                  .map(lambda word: (word, 1))  
                  .reduceByKey(lambda a, b: a + b)  
output = counts.collect()
```



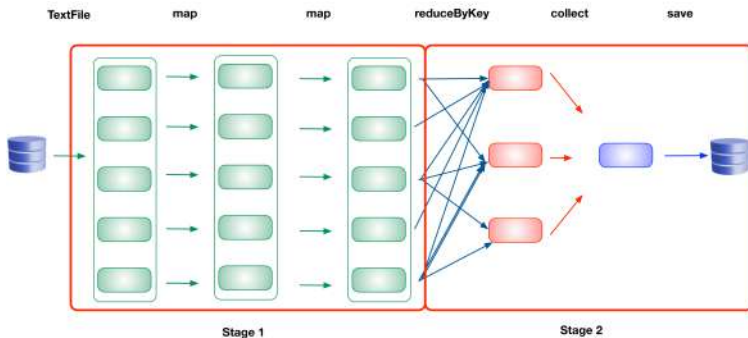
Example in Python

```
text_file = sc.textFile(''hdfs://... '')
counts = text_file.flatMap(lambda line: line.split(" "))
                    .map(lambda word: (word, 1))
                    .reduceByKey(lambda a, b: a + b)
output = counts.collect()
output.saveAsTextFile(''hdfs://... '')
```



Example in Python

2 Stages:





Stages

Stages are sequences of RDDs, that **don't have** a Shuffle in between

- Spark:
 - Creates a task for each Partition in the new RDD
 - Schedules and assigns tasks to slaves
 - All this happens **internally**



Summary

- **Task**: The fundamental unit of execution in Spark
- **Stage**: Set of Tasks that run in parallel
- **DAG**: Logical Graph of RDD operations
- **RDD**: Parallel dataset with partitions

References I



Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Michael Burrows, Tushar Chandra, Andrew Fikes, and Robert Gruber.

Bigtable: A distributed storage system for structured data (awarded best paper!).

In 7th Symposium on Operating Systems Design and Implementation (OSDI '06), November 6-8, Seattle, WA, USA, pages 205–218, 2006.



Holden Karau, Andy Konwinski, Patrick Wendell, and Matei Zaharia.

Learning Spark: Lightning-Fast Big Data Analytics.

O'Reilly Media, Inc., 1st edition, 2015.

References II



A growing number of applications are being built with Spark.
<https://www.oreilly.com/>.
Accessed 2016.



Spark Hadoop.
<http://spark.apache.org/docs/latest/cluster-overview.html>.
Accessed 2016.