

Lecture 2 – Virtualization and Containerization

Gianluca Quercini

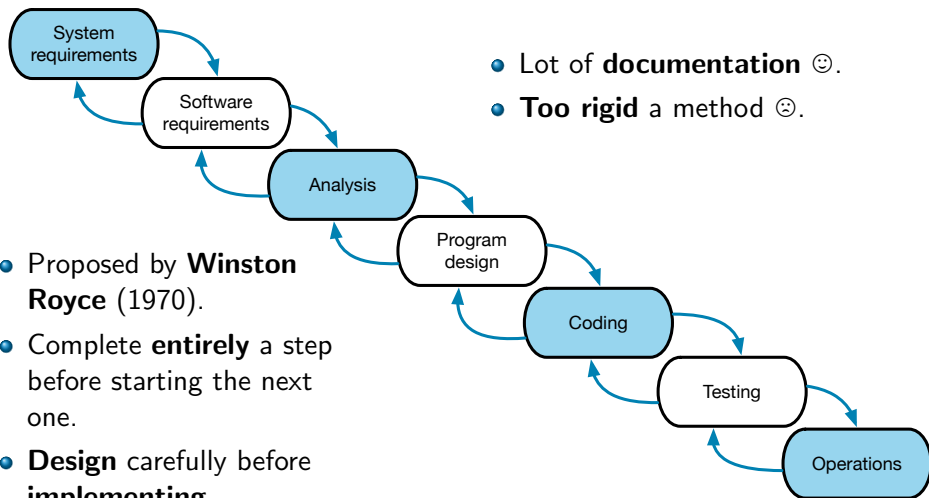
CentraleSupélec

SG8, 2019 – 2020

A bit of history


- Evolution of **software development methodologies**.
 - Waterfall (1970).
 - Agile (early 2000).
 - DevOps (2008-2009).
- Evolution of **software architectural styles**.
 - Monolithic.
 - Microservices.

Waterfall method



- Proposed by **Winston Royce** (1970).
- Complete **entirely** a step before starting the next one.
- **Design** carefully before **implementing**.

Towards Agile




- **RAD**: prototyping.
- **DSDM**: timeboxing, MoSCoW.
- **Scrum**: sprints, daily standups.
- **Extreme programming**: test-driven development.
- **Agile**: frequent software releases, feedback from the customers, embrace changes.

Towards DevOps

Developers

Operations



- **Developers** write some code.
- The **Quality Assessment (QA)** team tests the code.
 - They might send the code back to the developers for bug fixing.
- The **Operations** team receives the code ready for production.
 - They might send the code back to the developers if it doesn't work.

Towards DevOps

Even in companies that embraced Agile, teams worked in **silos**.

- **Developers** don't know the **production environment**.
- **Operations** don't know the **development environment**.
- When something doesn't work, they **blame each other**.

Developers and operations have **conflicting priorities**.

- **Developers** are evaluated on the number of **new features**.
- **Operations** are evaluated on the **system stability**.
- But changes threaten stability.



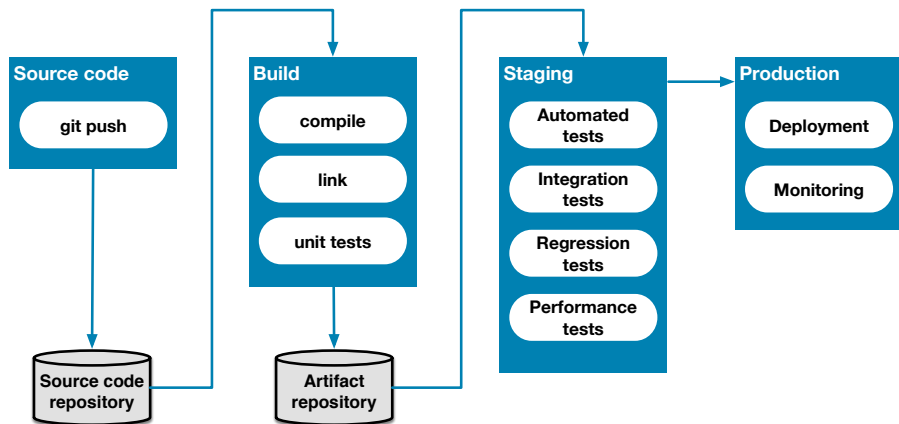
► [Image source](#)

DevOps

Cultural movement to address the Developers (**Dev**) - Operations (**Ops**) divide.

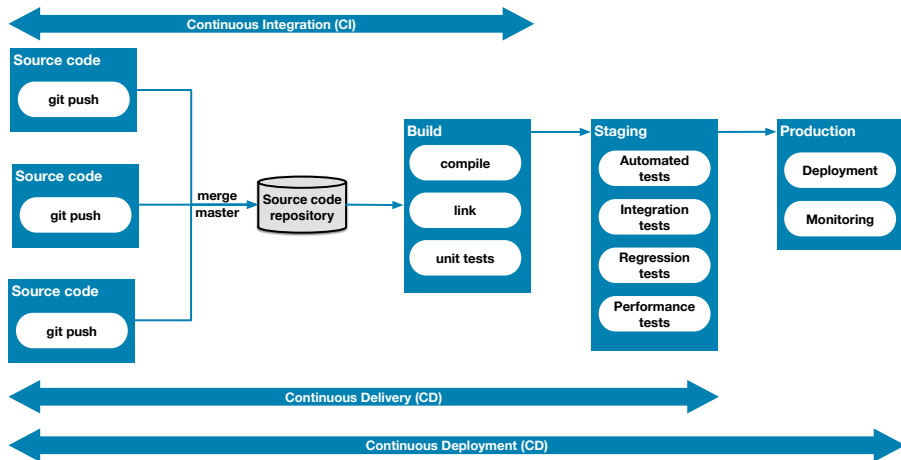
- **Collaborative learning** among the teams.
- **Automating** the software delivery pipeline.
- **Continuous feedback** in the teams.
 - Solve the problems as early as possible.
- **Continuous experimentation**, to learn from failures.
- **Collaboration** towards common goals.

DevOps — Software delivery pipeline



► Based on this description


DevOps — CI/CD



The key to DevOps practices is **automation**.

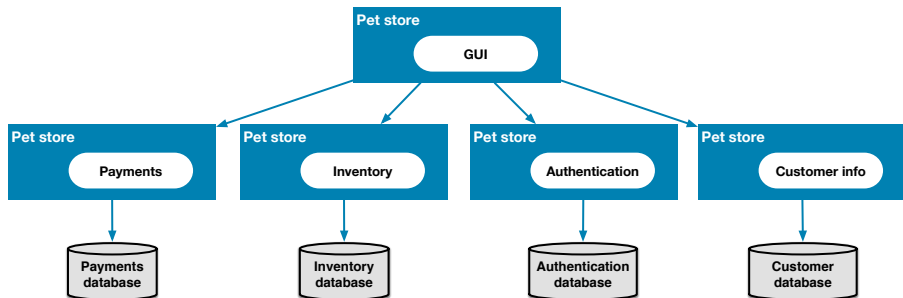
Monolithic architecture

- **Traditional applications** have a **monolithic architecture**.
- All features and services are coded into a single application.
- **Example.** Pet store application. [▶ Credit](#)



Microservices architecture

- Application composed of many **loosely coupled** and **independently deployable** smaller components, called **microservices**. [▶ Ref](#)
- Each microservice implements a specific feature of the application.
- Microservices interact by using APIs.



Microservices — Advantages


- **DevOps principles.**
 - Modularity.
 - Continuous delivery and deployment.
 - Failure isolation.
- Different microservices, **different technologies.**
- Microservices **scaled independently** of one another.
 - With a monolithic architecture, the whole application needs to be scaled even if only one module experiences an increased load.
- **Decentralized databases.**
 - Custom database technology for each microservice.
- Good architectural choice for **cloud-native applications.**

Microservices: package and deploy

How are microservices **packaged** and **deployed**? [▶ Ref](#)

- Multiple service instances per host.
- Single service instance per host.
- Service instance per **virtual machine**.
- Service instance per **container**.
- **Serverless** deployment.

Multiple service instances per host



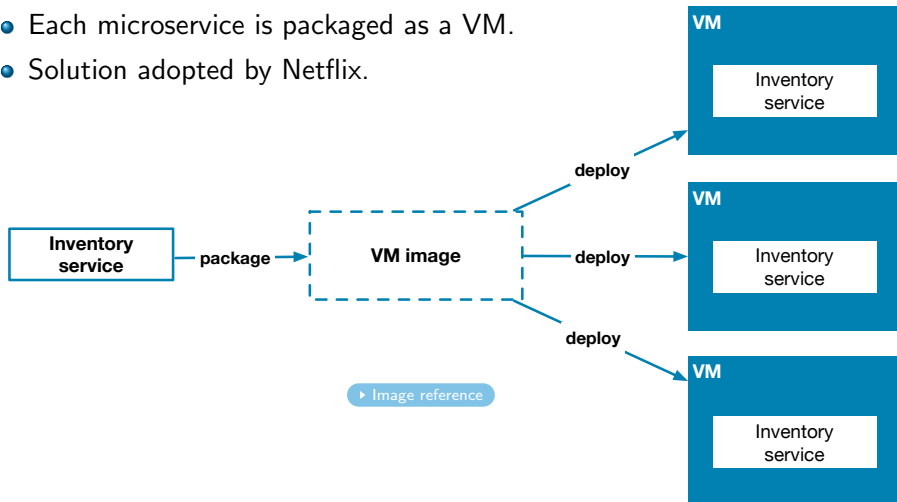
- Each host runs multiple service instances.
- **Efficient** use of resources 😊.
- **Limited isolation** of the service instances 😞.
- **Difficult to limit** the resources that a service uses 😞.
- **Difficult deployment** 😞.
 - Services are written with different languages and frameworks.

Single service instance per host

- Each single service instance is deployed on one host.
- **Isolation** between service instances 😊.
- **Easy monitoring** of the resources used by a service 😊.
- **Easy deployment** 😊.
- **Inefficient use of the resources** 😞.

Service instance per virtual machine (VM)

- Each microservice is packaged as a VM.
- Solution adopted by Netflix.

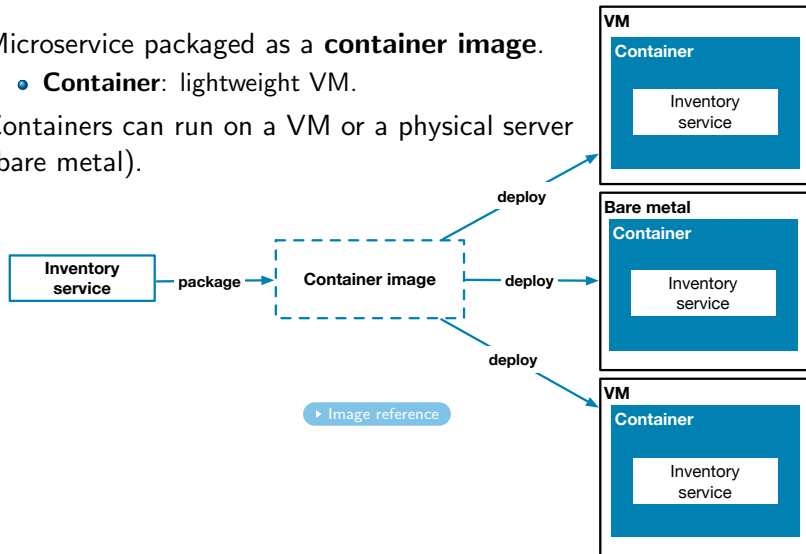


Service instance per virtual machine (VM)

- **Complete isolation** 😊.
- **Fixed amount of resources** (e.g., CPU and memory) per VM 😊.
- **Encapsulation.** Technological details hidden in the VM 😊.
 - All services are started and stopped in the same way.
- Mature **cloud infrastructure**: load balancing and autoscaling 😊.
- VMs images are **slow to build** and instantiate 😊.
 - A VM typically contains a full-blown operating system.
 - Solutions exist to generate lightweight images.

Service instance per container


- Microservice packaged as a **container image**.
 - **Container**: lightweight VM.
- Containers can run on a VM or a physical server (bare metal).



Service instance per container

- All advantages of VMs and
- Container images are **fast to build** 😊.
 - Containers don't encapsulate an entire operating system.
 - Containers only include the service and only what it takes to run the service.
- **Less mature** cloud infrastructure than for VMs 😞.

Serverless deployment



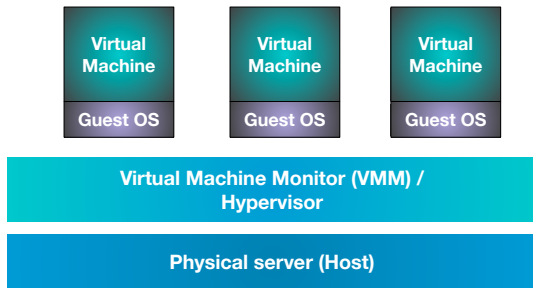
- Microservice packaged as an archive (e.g., zip, tar).
- The **serverless computing platform** runs as many instances of the service as needed.
- No visibility into the underlying infrastructure.

Serverless deployment

- **No need to configure** the underlying infrastructure 😊.
- **Automatically scale** the service to handle the load 😊.
- **Pay per request**, instead of paying for the infrastructure 😊.
- **Limited** language support 😞.
- **Unfit** for long-running services 😞.
- **No control** over the performances 😞.

Virtualization

- **Abstraction** of some physical resource into a logical object.
- Gerald Popek, Robert Goldberg (1974). First virtualization framework.
 - **Fidelity.** Virtual environment = host environment.
 - **Isolation.** Hypervisor has a complete control over system resources.
 - **Performance.** VM performances \approx Host performances.

[▶ Reference](#)

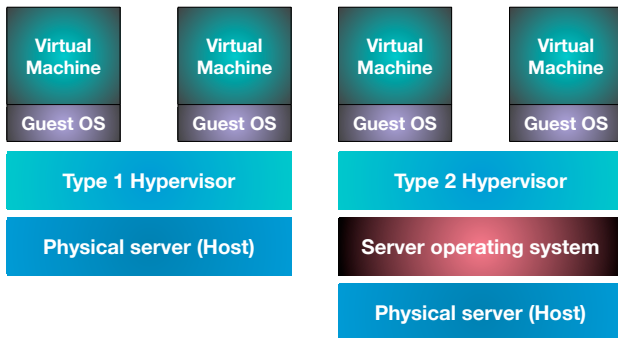
Virtualization

- Many factors pushed the adoption of the **one server, one application** policy.
 - Operating system (Microsoft Windows, as a **single-user system**).
 - Need to isolate applications.
 - Low hardware costs.
- Data centers grew larger and larger.
 - Difficult maintenance, energy consumption...
- **Inefficient use** of computing resources.
 - Many servers were idle 95% of the time.
- **Virtualization** allowed to stuff **several virtual servers** into a single physical server.
 - This is called **server consolidation**.
 - Better resource use, cost savings, less energy consumption.

[▶ Reference](#)

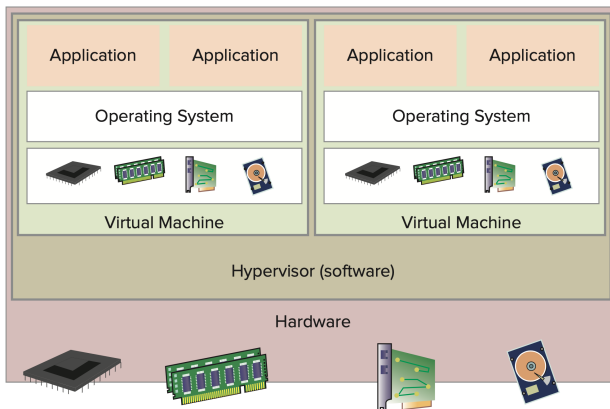
Hypervisor

- **Software component** that manages the interactions between the virtual machines and the underlying hardware.
- Type 1 (*bare metal*) hypervisors: VMware ESX, Microsoft Hyper-V.
- Type 2 hypervisors: VirtualBox.



Role of the Hypervisor

- Present an **abstract view** of the **physical resources** to the guests.
- Allocate the resources when they are requested from the guests.

[▶ Reference](#)

Virtual Machines

A **virtual machine** (VM) consists of the following components:

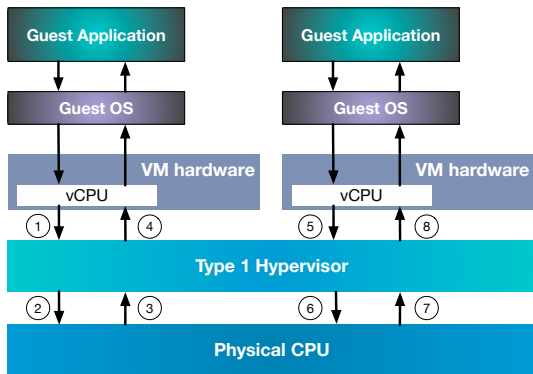
- One or more **applications**.
- An **operating system**.
- A set of **virtual hardware devices**.

On disk, a **virtual machine** is a **collection of files**:

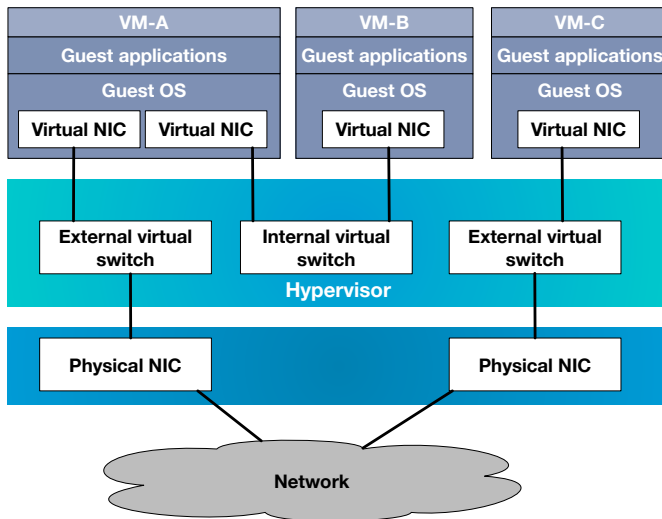
- Files that compose the virtual machine itself.
- A **configuration file**.
 - List of hardware devices: CPU, memory, storage.
- Virtual **disk files**.
 - Content of a virtual hard drive.

CPU in virtual machines

- The hypervisor assigns **time slices** on the physical CPU to execute instructions.
- No one-to-one mapping between virtual and physical CPUs.



Networks in virtual machines

[▶ Reference](#)

Networks in virtual machines

- **Virtual switches** play a similar role to physical switches.
 - Connect several devices to each other.
- Combination of **external** and **internal** virtual switches.
- **Secure** applications and servers.
 - VM-B cannot be accessed through the physical network.
 - Traffic between VM-A and VM-B is transparent to the physical network.
- **Fast communication** between virtual servers.
 - VM-A and VM-B communication happens **in memory**.
- Possibility of **migrating** a VM from a physical server to another.
 - Without interrupting the guest applications.

Advantages of virtual machines

- **Multiple virtual servers** in a physical server.
 - **Better** resource use.
 - **Lower** costs.
 - **Simplified** data center management.
- **Easy creation** of virtual machines from templates or clones.
 - A virtual server can be up and running in hours.
 - A physical server may need weeks to be up and running.
- **Easy migration** of virtual machines.
 - Reduced or minimized **downtime**.

Why containerization?

- Virtual machines: **good solution** to package a microservice.
 - The microservice is **embedded** into the VM.
 - **Easy deployment** on any machine.
- However, virtual machines are **heavy**.
 - They contain a **full-blown operating system** (e.g., Linux, Windows).
 - Starting or migrating a VM can be **time-consuming**.

Using a VM to package a single application is tantamount to use a “gigantic ship to transport a truck load of bananas”.


— Gabriel N. Schenker, [▶ Reference](#)

Why containerization?



Back in the old days, loading and unloading ships was difficult as each product came with a **different package**. [▶ Image source](#)


Then, there came **containers**, simple boxes with **standard dimensions**. [▶ Image source](#)



Why containerization?

- Each application came with a **“different package”**.
- Deployment was **application-dependent**.
- Difficult **dependency management** in production.

Developers



Operations


- **Containers** standardize application **development** and **deployment**.
- Great solution in a **cloud environment**.

What are containers?

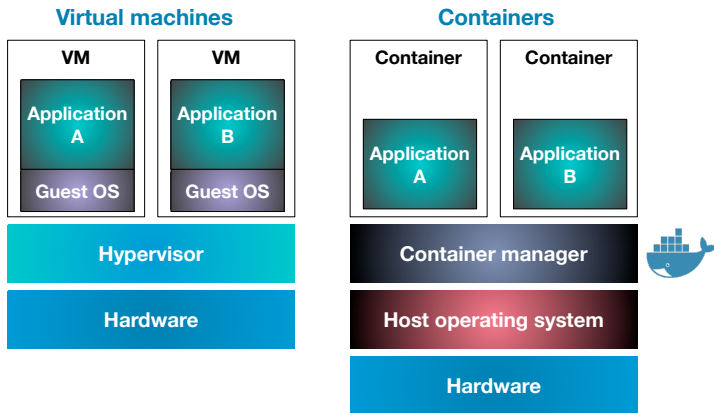
Definition (Container)

A **container** is standard unit of software that packages up code and all its dependencies so the application runs quickly and reliably from one computing environment to another. [▶ Reference](#)

- A **container** packages a piece of software with **all its dependencies**.
- Alternative (lighter) solution to virtual machines [▶ Reference](#).



Virtual machines and containers: differences



The container manager that we study in this course is **Docker**.

What is Docker?

Definition (Docker)

Docker is an **open platform** for **developing**, **shipping**, and **running** applications. It provides the ability to package and run an application in a loosely isolated environment called a **container** [▶ Reference](#).

Image

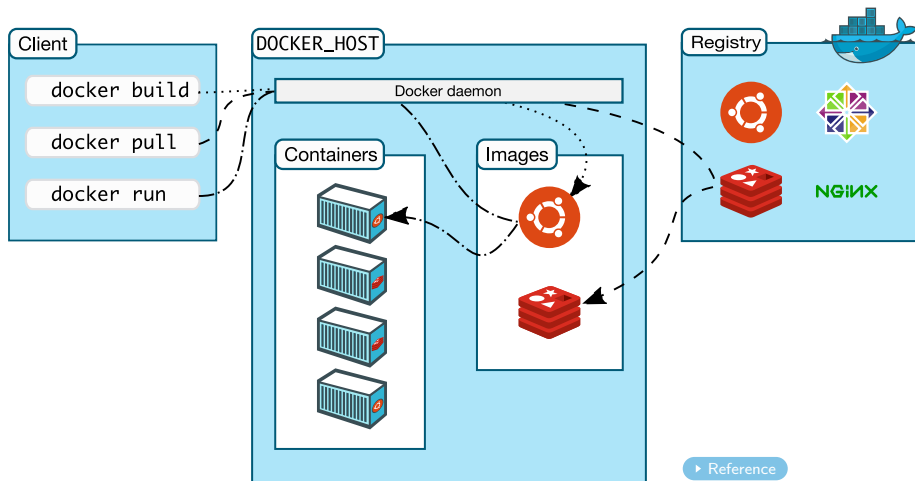
```
FROM python-3.7
```

```
RUN pip install requirements.txt
```

```
RUN python main.py
```


- Docker creates and runs a container from an **image**.
- **Image: read-only template** with instructions for creating and running a container.
- An image contains all files necessary to run an application in a container.

How does Docker work?

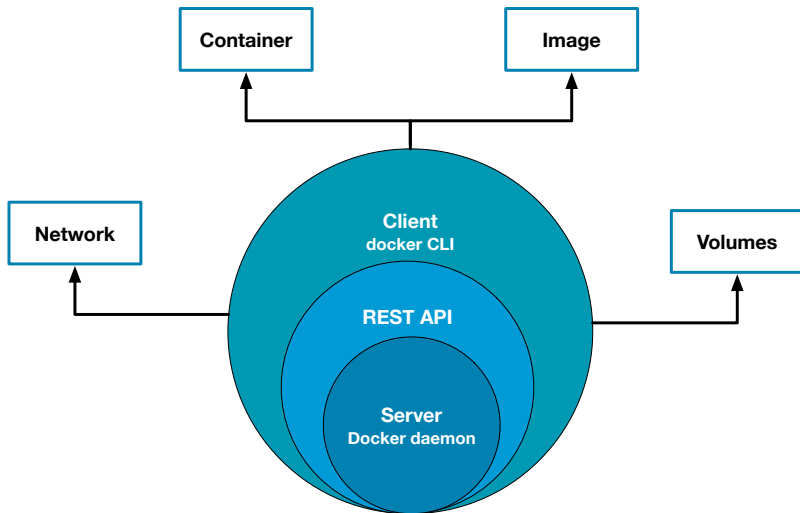


High-level Docker architecture

- **Docker engine.** Client-server application for building and containerizing applications.
- **containerd.** It manages the complete container lifecycle (image pull, container execution, low-level storage and network)
- **runc.** Command-line tool to run containers. Invoked by **containerd** with the proper parameters.
- Docker uses several features of the **Linux kernel**: **namespaces** (isolated workspaces), **control groups** (resource limitation), **union file systems** (build images).



Docker engine



Docker on MacOS

Older Mac systems

Docker Toolbox

Docker engine

Linux VM - boot2docker

VirtualBox

MacOS

Hardware

MacOS 10.13 or newer

Docker Desktop

Docker engine

Linux VM - LinuxKit

Hyperkit

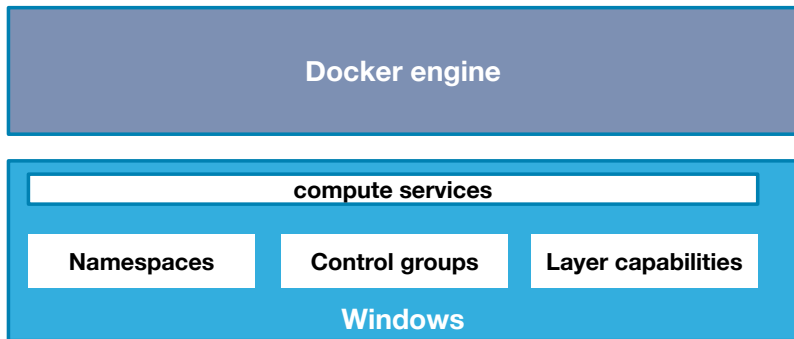
MacOS

**Hypervisor
framework**

Hardware


Docker on Windows

- Docker Engine **natively** supported in **Windows Server** (since 2016).
- The Windows Server kernel includes the technology to run containers.
- Definition of **Windows containers**.




Docker on Windows

Older Windows systems



Windows Server 2016+, Windows 10 Pro, Enterprise or Education



Running containers

```
docker run [options] image-name [command] [arg]
```

- **options**. List of options.
- **image-name**. Fully qualified name of the image used to run the container.
- **command**. The command to be executed in the container.
- **arg**. The arguments taken by the command executed in the container.
- The arguments between square brackets are **optional**.
 - The only **mandatory argument** is the **image name**.

Fully-qualified image names

registry.redhat.io/rhel8/mysql-80:1-69

- Name of the registry providing the image.
 - **Default:** registry.hub.docker.com
- Name of the user or organization owning the image.
 - **Default:** library
- Name of the image.
 - **Mandatory.**
- Tag (or, version) of the image.
 - **Default:** latest

Image names

Ubuntu image on DockerHub

- `registry.hub.docker.com/library/ubuntu:latest`
- `registry.hub.docker.com/library/ubuntu:18.04`
- `ubuntu:bionic`
- `ubuntu`

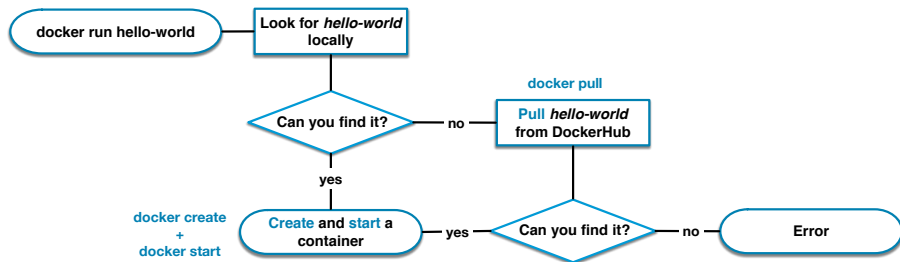
Supported tags and respective Dockerfile links

- `18.04`, `bionic-20200311`, `bionic`, `latest`
- `19.10`, `eoan-20200313`, `eoan`, `rolling`
- `20.04`, `focal-20200319`, `focal`, `devel`
- `14.04`, `trusty-20191217`, `trusty`
- `16.04`, `xenial-20200212`, `xenial`

[▶ Reference](#)

What happens when we run a container?

docker run hello-world



- The **containerized application** (the one inside the container) is run.
- When the containerized application terminates, the container is **stopped**.

How to list containers?

To list **running** containers:

```
docker container ls
```

To list **all** containers (running and stopped):

```
docker container ls -a
```

Output:

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
a997a41e815c	hello-world	"/hello"	6 seconds ago	Exited (0) 5 seconds ago		peaceful_roentgen

- The command executed when we run the container is *hello*.
- This is a **binary file** stored in the container.

Useful options of *docker run*

Give a container a name:

```
docker run --name my-hello-container hello-world
```

Remove a container when it stops:

```
docker run --rm hello-world
```

Interact with a container:

```
docker run -it ubuntu
```

- The last command opens a shell into the container.
- The shell is used to navigate the file system of the container.

Execute a command inside the container

Run `docker container ls -a` and look at the container created from the image *ubuntu*:

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
c03377554071	ubuntu	"/bin/bash"	11 minutes ago	Exited (0)		peaceful_fervent_cori

- When the container starts, the command `/bin/bash` is executed.
 - Bash is a Unix shell (command-line terminal).
- We can pass that shell any Linux command to execute.

```
docker run -rm ubuntu ls
```

What will this command output be?

Execute a command with arguments inside a container

- We pass any command executed inside a container some **arguments**.

The command *ping*

The command *ping* is used to test the reachability of a host on a IP network.

We run a container from the image *alpine* (yet another Linux distribution):

```
docker run --rm alpine ping www.centralesupelec.fr
```

- **ping** is the **command**.
- **www.centralesupelec.fr** is the **argument**.

Create, start and run: let's clarify

- The command **docker run** wraps two Docker commands:
 - **docker create** that creates a container.
 - **docker start** that starts the container.

Example

```
docker run -rm alpine ping www.centralesupelec.fr
```

is equivalent to :

```
docker create -rm --name my-ping alpine ping www.centralesupelec.fr
```

```
docker start -a my-ping
```

- The option **-a** attaches the **standard output** of the container.
- The container starts with the parameters set by **docker create**.
 - The command **ping www.centralesupelec.fr** is executed.
 - When the container is **stopped**, it gets **removed** (option **-rm**).

Create, start and run: an example

What do you obtain from the following command?

```
docker run -it --name my-ubuntu ubuntu
```

If you stop the container *my-ubuntu*, how would you start it again?

What is an image?


Definition (Image)

An **image** is a read-only template with instructions for creating a Docker container. [▶ Reference](#)

- An image contains a collection of **files** and **folders** necessary to run a containerized application.
- The files and folders are organized into a **layered filesystem**.

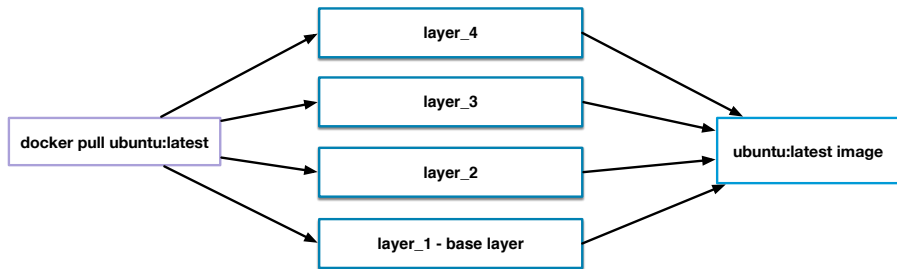
Anatomy of an image

- An image is a collection of **layers**.
- Each layer contains files and folders.
- Each layer has an **identifier**.
- Importantly, layers are **read-only**.



Pulling the layers

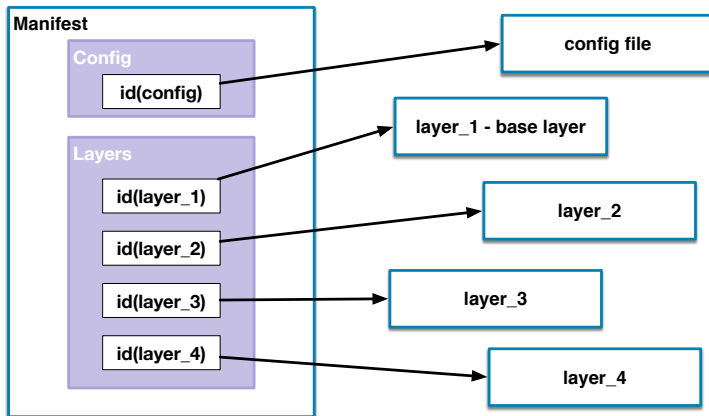
- Layers are **independent** of any image.
- To **pull an image** from the registry, Docker:
 - Pulls each layer **separately**.
 - Assembles the layers into a **layered filesystem**.



- How does Docker retrieve the layers? The **manifest file**.

The manifest file

- **ubuntu:latest** is associated to a file called the **manifest**.
 - Pointer to a **configuration file**: info to run a container from the image.
 - Pointer to the files containing the layers.



Manifest file of *ubuntu:latest*

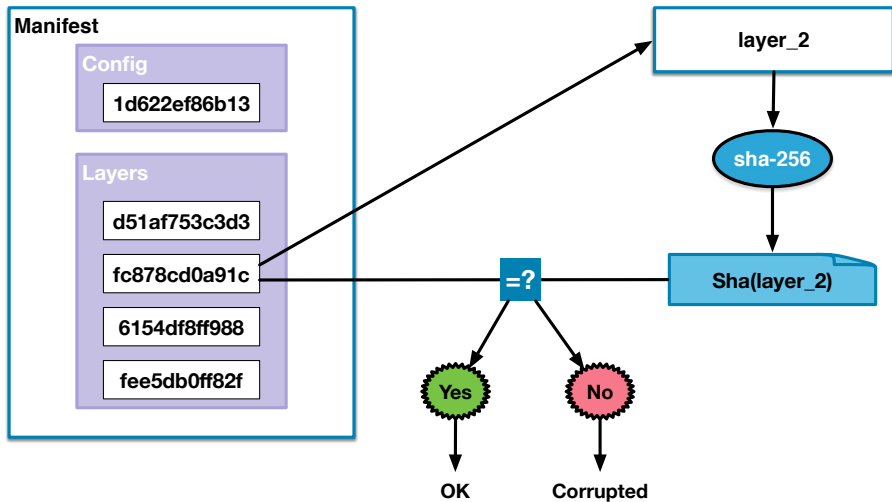
```
{
  "schemaVersion": 2,
  "mediaType": "application/vnd.docker.distribution.manifest.v2+json",
  "config": {
    "mediaType": "application/vnd.docker.container.image.v1+json",
    "size": 3408,
    "digest": "sha256:1d622ef86b138c7e96d4f797bf5e4baca3249f030c575b9337638594f2b63f01"
  },
  "layers": [
    {
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",
      "size": 28556247,
      "digest": "sha256:d51af753c3d3a984351448ec0f85ddafc580680fd6dfce9f4b09fdb367ee1e3e"
    },
    {
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",
      "size": 32304,
      "digest": "sha256:fc878cd0a91c7bece56f668b2c79a19d94dd5471dae41fe5a7e14b4ae65251f6"
    },
    {
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",
      "size": 847,
      "digest": "sha256:6154df8ff9882934dc5b5f265b8b85a3aeada06387447ffa440f7af7f32b0e1d"
    },
    {
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",
      "size": 163,
      "digest": "sha256:fee5db0ff82f7aa5ace63497df4802bbad8f2779ed3e1858605b791dc449425"
    }
  ]
}
```

Layer identifiers

- **Identifier** of a layer: **SHA256 digest** of its content.
- **Identifier** of the configuration file: a **SHA256 digest** of its content.
- **Digest** of a file: cryptographic representation of the file **content**.
 - Computed with the **SHA-256 algorithm**.
- Important to verify the **integrity** of a layer.
- Two files have the **same digest** iff they have the **same content**.

$$Sha(x) = Sha(y) \iff x = y$$

Layer integrity



Sharing layers

- Two images might reference the same layers.
- Layers are only downloaded once.



Example: pulling *ubuntu:latest*

```
docker pull ubuntu
```

Using default tag: latest

latest: Pulling from library/ubuntu

d51af753c3d3: Pull complete

fc878cd0a91c: Pull complete

6154df8ff988: Pull complete

fee5db0ff82f: Pull complete



Pulling each layer independently

Digest:

sha256:747d2dbbaaee995098c9792d99bd333c6783ce56150d1b11e333bbceed5c54d7

SHA-256 digest of the manifest file of the image

Example: pulling quercinigia/ubuntu-ping:latest

```
docker pull quercinigia/ubuntu-ping
```

Using default tag: latest

latest: Pulling from quercinigia/ubuntu-ping:latest

d51af753c3d3: Already exists

fc878cd0a91c: Already exists

6154df8ff988: Already exists

fee5db0ff82f: Already exists

2d2158c9757f: Pull complete



Layers from ubuntu:latest

Digest:

sha256:0ae0d96e80f456f7d078b2cff5983ee26a0767260fc5c4df11141b224c74ff0a

Image identifier

Command to list local images:

```
docker images
```

Or:

```
docker image ls
```

The output is:

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
ubuntu	latest	1d622ef86b13	3 days ago	73.9MB

SHA-256 digest of the
configuration file of the image

How to build an image

- **Working example:** sentiment analysis of sentences in Python.

Source main.py

```
import sys
import nltk
from nltk.sentiment.vader import SentimentIntensityAnalyzer

if __name__=='__main__':
    nltk.download('vader_lexicon')
    sid = SentimentIntensityAnalyzer()
    sentence = sys.argv[1]
    print("\nSentiment analysis on the following sentence ", sentence)
    print(sid.polarity_scores(sentence))
```

- The image must include:
 - The source *main.py*
 - All the dependencies (e.g., *nltk*).
 - The Python interpreter.
 - Instructions on how to run the code.
- Most common way to build an image: **Dockerfile**.

Dockerfile

- **Each line** of the Dockerfile corresponds to a **layer** in the image.
- **FROM**. Specifies the **base image**.
- **RUN**. Used to run any valid Linux command.
- **WORKDIR**. Sets the working directory in the image.
- **ENTRYPOINT**. Command to execute when a container is run.
- **CMD**. Any argument passed to the entrypoint command.

Build instructions	{	FROM python:3.7-slim	build from base image
		RUN mkdir -p /app	create folder /app in the image
		WORKDIR /app	set working directory to /app
		COPY ./main.py ./requirements.txt /app/	copy files to /app
Run instructions	{	RUN pip install -r requirements.txt	install the Python dependencies
		ENTRYPOINT ["python", "main.py"]	run python main.py...
		CMD ["A default sentence"]with an argument

Build the image

```
docker build -t sentiment-analysis:latest .
```

- The Docker engine looks for a file named **Dockerfile** in the current directory.
- Executes the instructions in the Dockerfile and creates an image tagged *sentiment-analysis:latest*.
- If the Dockerfile has another name (e.g., *Dockerfile-sentiment*), the command to build the image is as follows:

```
docker build -t sentiment-analysis:latest -f Dockerfile-sentiment .
```

The resulting image

sentiment-analysis:latest

CMD ["A default sentence"]

ENTRYPOINT ["python", "main.py"]

RUN pip install -r requirements.txt

COPY ./main.py ./requirements.txt /app/

WORKDIR /app

RUN mkdir -p /app

FROM python_3.7-slim

Image filesystem

nlTK library files

**files main.py,
requirements.txt**

Directory /app

Python installation files

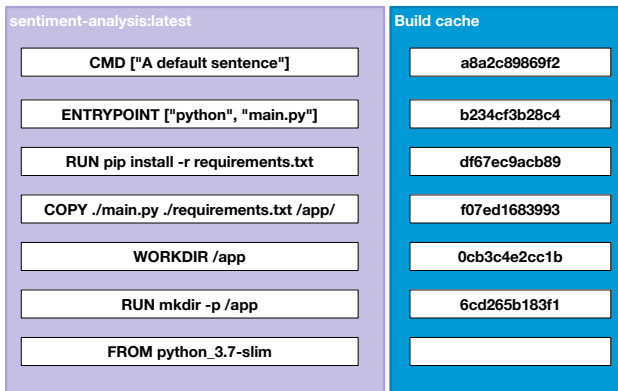
A closer look at the image history

docker history sentiment-analysis

IMAGE	CREATED	CREATED BY	SIZE	COMMENT
a8a2c89869f2	20 seconds ago	/bin/sh -c #(nop) CMD ["A default sentence"]	0B	
b234cf3b28c4	20 seconds ago	/bin/sh -c #(nop) ENTRYPOINT ["python" "mai...	0B	
df67ec9acb89	21 seconds ago	/bin/sh -c pip install -r requirements.txt	18.6MB	
f07ed1683993	26 seconds ago	/bin/sh -c #(nop) COPY multi:4e5c15fb4242a23...	334B	
0cb3c4e2cc1b	26 seconds ago	/bin/sh -c #(nop) WORKDIR /app	0B	
6cd265b183f1	26 seconds ago	/bin/sh -c mkdir -p /app	0B	
59ade49ea505	3 days ago	/bin/sh -c #(nop) CMD ["python3"]	0B	
<missing>	3 days ago	/bin/sh -c set -ex; savedAptMark="\$(apt-ma...	7.51MB	
<missing>	3 days ago	/bin/sh -c #(nop) ENV PYTHON_GET_PIP_SHA256...	0B	
<missing>	3 days ago	/bin/sh -c #(nop) ENV PYTHON_GET_PIP_URL=ht...	0B	
<missing>	3 days ago	/bin/sh -c #(nop) ENV PYTHON_PIP_VERSION=20...	0B	
<missing>	3 days ago	/bin/sh -c cd /usr/local/bin && ln -s idle3...	32B	
<missing>	3 days ago	/bin/sh -c set -ex && savedAptMark="\$(apt-...	95.1MB	
<missing>	4 days ago	/bin/sh -c #(nop) ENV PYTHON_VERSION=3.7.7	0B	
<missing>	4 days ago	/bin/sh -c #(nop) ENV GPG_KEY=0D96DF4D4110E...	0B	
<missing>	4 days ago	/bin/sh -c apt-get update && apt-get install...	7.03MB	
<missing>	4 days ago	/bin/sh -c #(nop) ENV LANG=C.UTF-8	0B	
<missing>	4 days ago	/bin/sh -c #(nop) ENV PATH=/usr/local/bin:/...	0B	
<missing>	4 days ago	/bin/sh -c #(nop) CMD ["bash"]	0B	
<missing>	4 days ago	/bin/sh -c #(nop) ADD file:9b8be2b52ee0fa31d...	69.2MB	

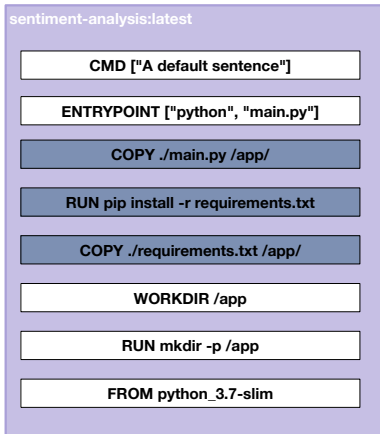
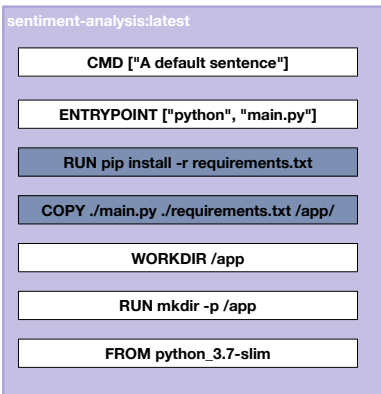
Build cache

- For each layer, the Docker engine creates a **cache image**.
- When building again, only the modified layer and those above are built again (**build cache**).



Build cache

- For efficient builds, order matters.
- When we modify the source code, the dependencies are installed again.



Dockerfile best practices

- **Small number of layers.**
 - Example: merge two consecutive RUN commands.
- **Order layers correctly.**
 - Efficient use of the build cache.
- **Small size images.**
 - Put only necessary files into images.

Running our container

What's the output of the following command?

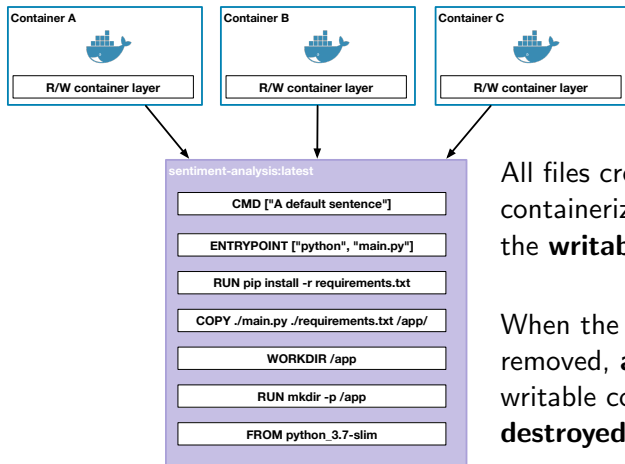
```
docker run -rm sentiment-analysis
```

What's the output of the following command?

```
docker run -rm sentiment-analysis "This is my positive sentence"
```


The container layer

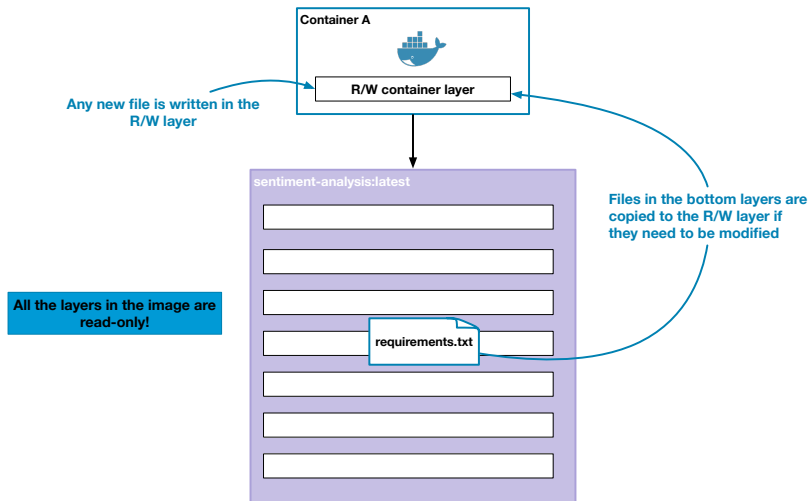
- A container is a **runnable instance** of an image.
- It includes all the layers of the image, and a **writable container layer**.



All files created by the containerized app are stored in the **writable container layer**.

When the container is removed, **all files** in the writable container layer are **destroyed**.

Running a container: copy-on-write (COW)



Highlights

- An image is composed of a **layered filesystem**.
- Layers are **independent** of the image.
 - Layers can be shared by images.
- Layers are **identified** by their **content digest**.
 - Possibility of **checking their integrity**.
- Layers are **immutable**.
 - Images, being a collection of immutable layers, are immutable too.
 - An image always behaves the same.
- A container is a **runnable instance** of an image with a **R/W top layer**.
- New files are written to the R/W container layer.
- Files in the bottom layers are copied to the R/W top layer before any modification.

Volumes

- All data created by a containerized application is written to the R/W container layer.
- The R/W layer is **destroyed** when the container is removed.

How can we persist data beyond a container's lifecycle?

- The answer is: **volumes**.

Volumes

Definition (Volume)

A **volume** is a directory that lives outside of the image filesystem and corresponds to a directory in the host computer.

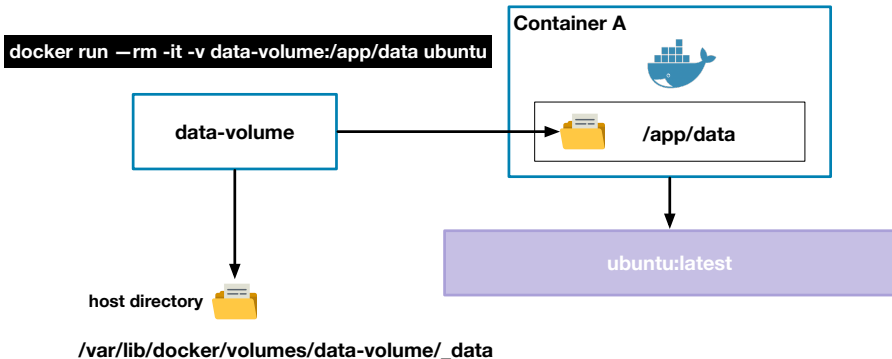
- We can think of a **volume** as an external storage drive for a container.
- Before using a volume, we need to declare it.

```
docker volume create data-volume
```

- *data-volume* is attached to a directory in the **host computer**.
 - Docker for Linux. Directory accessible under `/var/lib/docker/volumes`.
 - Docker for Mac. Directory accessible in the Linux VM.
 - Docker for Windows. Directory accessible in the Hyper-V container.


Mounting a volume

- All files written in the directory `/app/data` are written to the volume `data-volume`.



Data sharing

- Volumes are a great way to **share data** between containers.



- The two containers don't need to be run from the same image.
- The two containers can mount the volume at any directory.

Networks

- Containers **isolate** the applications from their environment.
- Great feature in a micro-service architecture.

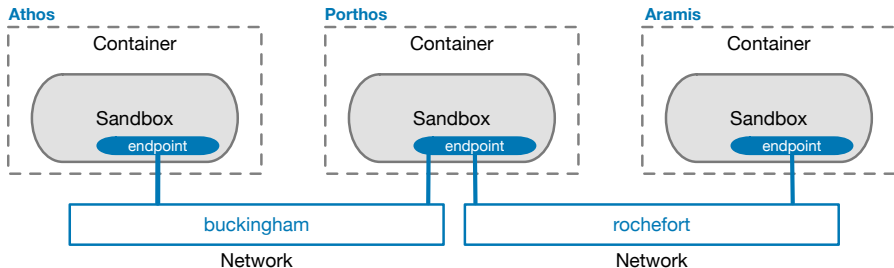
How can two containers running on the same host **communicate**?

- Docker provides an object called **network**.

Container Network Model

Docker defines a **container network model** consisting of three elements:

- **Sandbox.** Contains the **network configuration** (IP and MAC addresses, routing tables, DNS records).
- **Endpoint.** Connection between the sandbox and a network.
- **Network.** Provides the functions to connect two or more endpoints.
 - A network implementation is called a **driver**.



Driver

Running the command `docker network ls` gives the following output:

NETWORK ID	NAME	DRIVER	SCOPE
7ca6f4877494	bridge	bridge	local
4140d2566476	host	host	local
b35755c21cd2	none	null	local

- **bridge**. Network based on Linux bridges.
- **host**. Network of the host.
- **null**. Used to disconnect containers from any network.

All these drivers allow **single-host networking**.

Creating networks

- Command to create a network:

```
docker network create buckingham
```

- Run a container and attach it to a network:

```
docker run -rm -it -name c1 --network buckingham alpine  
docker run -rm -it -name c2 --network buckingham alpine
```

- If we need to attach a container to another network:

```
docker network connect bridge c1
```

Inspecting a network


docker inspect buckingham

```
....  
"IPAM": {  
  "Driver": "default",  
  "Options": null,  
  "Config": [  
    {  
      "Subnet": "172.18.0.0/16",  
      "Gateway": "172.18.0.1"  
    }  
  ]  
},  
....
```

```
....  
"Containers": {  
  "4ba0f09c5f7b": {  
    "Name": "c2",  
    "EndpointID": "5557c6fe08f2",  
    "MacAddress": "02:42:ac:12:00:03",  
    "IPv4Address": "172.18.0.3/16",  
    "IPv6Address": ""  
  },  
  "b971823f57c0": {  
    "Name": "c1",  
    "EndpointID": "5fbaccf81b42",  
    "MacAddress": "02:42:ac:12:00:02",  
    "IPv4Address": "172.18.0.2/16",  
    "IPv6Address": ""  
  }  
},  
....
```


Inspecting a network

- Docker daemon acts as a **DHCP server**.
- The **hostname** of a container is **its identifier**.
- But one can also use its name.



► Image source

Publishing ports



```
docker run --rm -it --name c1 --network buckingham -p 8080:7070 alpine
```

Resources and implementation

- A docker primer [▶ Click here](#) .
- Tutorial: Getting started with Docker [▶ Click here](#) .



[▶ Image source](#)